

Sifei Luan

EDUCATION

| | |
|---|-----------------------------|
| University of California, Berkeley <i>Ph.D. in Computer Science</i> Research topics: AI & ML systems, distributed data processing systems, cloud computing | Berkeley, CA 2020 – 2024 |
| The University of Chicago <i>B.S. in Computer Science and B.A. in Statistics</i> | Chicago, IL 2013 – 2017 |

HONORS & AWARDS

| | |
|---|------|
| CloudSort World Record <i>The Sort Benchmark Committee</i> A foundational big data benchmark. The most cost-efficient way to sort 100 TB of data a public cloud [1, 2]. | 2022 |
| IEEE Software Magazine Best Paper Award <i>IEEE Computer Society</i> Awarded to the article “AI in Software Engineering at Facebook” [5]. | 2021 |
| ACM SIGPLAN Distinguished Paper Award <i>The ACM Special Interest Group on Programming Languages</i> Awarded to the paper “Aroma: Code Recommendation via Structural Code Search” [8]. | 2019 |
| Champion of the Midwest Trading Competition <i>The University of Chicago</i> Developed the best-performing automated trading algorithms for three securities markets. | 2016 |
| Regional Finalist of the ACM International Collegiate Programming Contest <i>Association for Computing Machinery</i> Mid-central USA Regional Contest Finalist. | 2014 |

EXPERIENCE

| | |
|---|----------------------------------|
| Wayo <i>Co-founder & CTO</i> ◦ Created an AI-powered global product sourcing company and raised \$2M in venture capital funding. ◦ Built and led a team of 10 engineers, designers and product managers to launch a marketplace with \$1M+ GMV. | San Francisco, CA 2023 – 2024 |
| Anyscale <i>Software Engineering Intern</i> ◦ Integrated Exoshuffle, a research system, into Ray Data, an open-source data processing framework for ML. | San Francisco, CA 2021 |
| Facebook <i>Software Engineer</i> ◦ Created ML-based software engineering tools for natural language code search, code recommendation, code completion, code synthesis, code review, and automatic bug fixing. ◦ Deployed these tools on codebases of 10B+ LOC, improving the productivity of 50K+ engineers. ◦ Published 3 academic papers and presented at top-tier academic conferences for software engineering. | Menlo Park, CA 2017 – 2019 |
| Facebook <i>Software Engineering Intern</i> ◦ Implemented several new features in the Facebook Pages SMB Platform, totaling 20K+ LOC. | Menlo Park, CA 2016 |
| SketchMe <i>Co-founder & CTO</i> ◦ Created a social app company and raised \$1.5M in venture capital funding. ◦ Built a team of 5 engineers and designers to launch a mobile app with 10K+ users. | Menlo Park, CA 2014 – 2015 |

PUBLICATIONS

- [1] Frank Sifei Luan, Stephanie Wang, Samyukta Yagati, Sean Kim, Kenneth Lien, Isaac Ong, Tony Hong, SangBin Cho, Eric Liang, and Ion Stoica. *Exoshuffle-CloudSort*. 2023. arXiv: 2301.03734 [cs.DC].
- [2] Frank Sifei Luan, Stephanie Wang, Samyukta Yagati, Sean Kim, Kenneth Lien, Isaac Ong, Tony Hong, Sangbin Cho, Eric Liang, and Ion Stoica. “Exoshuffle: An Extensible Shuffle Architecture”. In: *Proceedings of the ACM SIGCOMM 2023 Conference*. ACM SIGCOMM ’23. New York, NY, USA: Association for Computing Machinery, 2023, pp. 564–577.
- [3] Zongheng Yang, Zhanghao Wu, Michael Luo, Wei-Lin Chiang, Romil Bhardwaj, Woosuk Kwon, Siyuan Zhuang, Frank Sifei Luan, Gautam Mittal, Scott Shenker, and Ion Stoica. “SkyPilot: An Intercloud Broker for Sky Computing”. In: *20th USENIX Symposium on Networked Systems Design and Implementation (NSDI 23)*. Boston, MA: USENIX Association, Apr. 2023, pp. 437–455.
- [4] Zongheng Yang, Wei-Lin Chiang, Sifei Luan, Gautam Mittal, Michael Luo, and Ion Stoica. “Balsa: Learning a Query Optimizer Without Expert Demonstrations”. In: *Proceedings of the 2022 International Conference on Management of Data*. SIGMOD ’22. Philadelphia, PA, USA: Association for Computing Machinery, 2022, pp. 931–944.
- [5] Johannes Bader, Sonia Seohyun Kim, Frank Sifei Luan, Satish Chandra, and Erik Meijer. “AI in Software Engineering at Facebook”. In: *IEEE Software* 38.4 (2021), pp. 52–61.
- [6] Stephanie Wang, Eric Liang, Edward Oakes, Ben Hindman, Frank Sifei Luan, Audrey Cheng, and Ion Stoica. “Ownership: A Distributed Futures System for Fine-Grained Tasks”. In: *18th USENIX Symposium on Networked Systems Design and Implementation (NSDI 21)*. USENIX Association, 2021, pp. 671–686.
- [7] Zongheng Yang, Amog Kamsetty, Sifei Luan, Eric Liang, Yan Duan, Xi Chen, and Ion Stoica. “NeuroCard: one cardinality estimator for all tables”. In: *Proc. VLDB Endow.* 14.1 (2020), pp. 61–73.
- [8] Sifei Luan, Di Yang, Celeste Barnaby, Koushik Sen, and Satish Chandra. “Aroma: code recommendation via structural code search”. In: *Proc. ACM Program. Lang.* 3.OOPSLA (Oct. 2019).
- [9] Saksham Sachdev, Hongyu Li, Sifei Luan, Seohyun Kim, Koushik Sen, and Satish Chandra. “Retrieval on source code: a neural code search”. In: *Proceedings of the 2nd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages*. MAPL 2018. Philadelphia, PA, USA: Association for Computing Machinery, 2018, pp. 31–41.

INVITED TALKS

| | |
|---|-------------------|
| Ray Data: Unlocking Heterogeneous Execution with the Streaming Batch Model | 2024 |
| <i>Paul G. Allen School of Computer Science & Engineering, University of Washington</i> | Seattle, WA |
| Exoshuffle: An Extensible Shuffle Architecture | 2024 |
| <i>Meta</i> | Menlo Park, CA |
| The Ray Dataplane: History and the CloudSort World Record | 2023 |
| <i>Ray Summit [video]</i> | San Francisco, CA |
| Exoshuffle: Large-Scale Shuffle at the Application Level | 2022 |
| <i>Google</i> | Mountain View, CA |
| Using ML for Code Discovery at Facebook | 2019 |
| <i>Curry On London (co-located with ECOOP) [video]</i> | London, UK |
| Using Machine Learning for Developer Productivity | 2019 |
| <i>F8 Developer Conference [video]</i> | San Jose, CA |

ACADEMIC SERVICES

| | |
|--|------|
| Reviewer | 2024 |
| <i>Conference on Neural Information Processing Systems (NeurIPS)</i> | |
| Reviewer | 2023 |
| <i>Conference on Neural Information Processing Systems (NeurIPS)</i> | |

| | |
|--|------|
| Artifact Evaluation Committee <i>The European Conference on Computer Systems (EuroSys)</i> | 2022 |
| Artifact Evaluation Committee <i>Symposium on Operating Systems Principles (SOSP)</i> | 2021 |
| Reviewer <i>Conference on Machine Learning and Systems (MLSys)</i> | 2021 |